

# **TDT4310, Lab session 2**

## **Back to basics: Tokenization, language modeling and word representations**

Tollef Jørgensen, January 30, 2024.

# Outline

- Tokenization
- Some basics of language modeling
- Word representations

# The assignments and blackboard


- Information about the course and assignments:
  - **Please mainly use the blackboard forum**
    - Can post anonymously as well
- Sensitive information/other questions:
  - Emails (**please add [TDT4310] in the subject**)
    - If related to an assignment (sick leave, etc.)
      - Me
    - Else:
      - Björn/me + CC TA's

# Reference group

- Pls
- Probably only 2 meetings total! There will be cookies
- Talk to me or email [gamback@ntnu.no](mailto:gamback@ntnu.no) / [tollef.jorgensen@ntnu.no](mailto:tollef.jorgensen@ntnu.no) :-)

# LLM usage detection for Lab 1

Distribution of answer scores



# LLM usage detection for Lab 1

- Approach: gather outputs from 6 models for the questions given in the Lab, then rank each answer from students on the following:
  - Sentence embeddings from BERT models to create a single vector for sentence representations (**related to this lab!**)
    - Compute cosine similarity between the LLM answer and your answers
    - Weighted by Jaccard similarity (set of overlapping words)
- Results: none seemed to rely solely on LLMs, although some answers were definitely *inspired*

# Tokenization

- The task of splitting a text into smaller pieces (tokens)
- Can be words, characters, subwords, ...
- Some methods
  - Whitespace
  - Regex (e.g., match URLs, hashtags)
  - Rules (such as entities, e.g. "New York")
  - Contractions (don't/do n't/don 't/don ' t)
  - Expressions (multi-word exp. tokenizer)
  - ... There's a whole lot

# Tokenization

- But limiting tokenization to just words is also problematic.
- What happens when you observe an out-of-domain token?
  - Your system should be able to handle *any* input



# Tokenization

- Sentence: "cats and rats"
- Byte pair encoding
  - Merge from tokenization to commonly observed subword units
    - Better suited for large corpora
  - Add frequent pairs (for a certain amount of iterations)
    - c+a: 1, a+t: 2, t+s: 2, ... => c + ats, and, r + ats
- WordPiece (used to pre-train BERT)
  - gather vocabulary (characters)
  - split on units and add a special ## prefix for joined tokens
  - "cats" -> "c", "##a", "##t", "##s". Calculate the frequency of each pair
    - here the frequency of ##t and ##s would be 2, and receive a higher score
      - merge and repeat: "c" "##a" "ts"

# Tokenization

- Going back to LLMs and *context length*
  - *number of tokens in context*
- If a word is commonly used, such as "cat", this is likely one token.
  - But what about other languages and rarely occurring strings?
    - A single japanese character (such as cat, 猫) counts as 3 tokens!

GPT-3.5 &amp; GPT-4

GPT-3 (Legacy)

Trondheim, Ålesund, Gjøvik

Clear

Show example

Tokens

11

Characters

26

Trondheim, Ålesund, Gjøvik

# Tokenization - Implementations in NLTK

- <https://www.nltk.org/api/nltk.tokenize>
- Chapter 3: <https://www.nltk.org/book/ch03.html>

# Language modeling

- Not as fancy as what we explored the first time around...
- Focused on n-grams and frequency distributions
  - bigrams, trigrams, ...
- The point is to get you to think about how can we use these simple methods to model language

# Language modeling

- Using co-occurrences in text will always have limitations
  - No context
  - Ambiguous words
    - running a marathon
    - the ink is running
    - she's running the show
- Given what you will hopefully learn about word representations, we can learn how these words appear together!

# Word Representations

- Much higher flexibility than working with tokens of words directly
- Encodes some form of representation for each word
- Simple examples:
  - One-hot encoding
  - Bag-of-words (BoW)
    - Mapping of word:count
  - The above are *sparse* representations. They don't scale well with infinitely many words!

```
word_vector_map = {  
    "cat": [0.1, 0.2, 0.3, 0.4],  
    "dog": [0.2, 0.3, 0.4, 0.5],  
    "mouse": [0.3, 0.4, 0.5, 0.6],  
    ...  
}
```


$$vec_{cat} = [1, 0, 0]$$

$$vec_{dog} = [0, 1, 0]$$


$$vec_{mouse} = [0, 0, 1]$$

# Word Representations

- Distributional hypothesis
  - *You shall know a word by the company it keeps (Firth, 1957)*
- Heavily represented in word embeddings
  - E.g., Word2Vec
- Two separate objectives (word2vec)
  - Continuous bag of words
    - Given a context, predict a word
  - Skip-gram
    - Given a word, predict context
- Suggested: <https://arxiv.org/pdf/1301.3781.pdf>



**CBOW**



**Skip-gram**